

---

# Exponential Smoothing

BSAD 8310: Business Forecasting — Lecture 3

---

Department of Economics  
University of Nebraska at Omaha  
Spring 2026

---

1 Motivation: Adaptive Averaging

2 Simple Exponential Smoothing

3 Holt's Linear Method

4 Holt-Winters Seasonal Method

5 The ETS Framework

6 Key Takeaways and Roadmap

---

## Motivation: Adaptive Averaging

Regression treats all observations equally — but recent observations carry more information.

---

---

In OLS regression, *every* observation receives equal implicit weight:

$$\hat{\beta} = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y} \quad \Rightarrow \quad \hat{y}_{T+h|T} = \mathbf{x}'_{T+h}\hat{\beta}$$

When the data-generating process **drifts over time** — structural breaks, changing growth rates, evolving seasonality — old observations **contaminate** the forecast.

A retailer's sales pattern shifts after a store renovation. OLS keeps fitting to pre-renovation data, systematically over- or under-forecasting for months afterward.

---

**Idea:** assign geometrically declining weights to past observations.

Observation	Age	Weight
$y_T$	0	$\alpha$
$y_{T-1}$	1	$\alpha(1 - \alpha)$
$y_{T-2}$	2	$\alpha(1 - \alpha)^2$
$y_{T-j}$	$j$	$\alpha(1 - \alpha)^j$

Weights sum to 1:  $\sum_{j=0}^{\infty} \alpha(1 - \alpha)^j = 1$  for  $0 < \alpha \leq 1$ .

A single parameter  $\alpha$  controls how fast old observations are **forgotten**.     $\alpha \approx 1$ : almost all weight on the latest value.     $\alpha \approx 0$ : long memory.

---

## Simple Exponential Smoothing

One parameter, one component: the level.

---

Let  $\ell_t$  denote the **level** at time  $t$ . SES updates:

$$\ell_t = \alpha y_t + (1 - \alpha) \ell_{t-1}, \quad 0 < \alpha \leq 1$$

with initialization  $\ell_0$  estimated from the data.

**Recursive expansion** reveals the exponential weights:

$$\ell_t = \alpha \sum_{j=0}^{t-1} (1 - \alpha)^j y_{t-j} + (1 - \alpha)^t \ell_0$$

As  $t \rightarrow \infty$ , the initialization term  $(1 - \alpha)^t \ell_0 \rightarrow 0$  for any  $\alpha > 0$  — the starting value eventually becomes irrelevant.

SES can be rewritten as an **error-correction equation**:

$$\ell_t = \ell_{t-1} + \alpha \underbrace{(y_t - \ell_{t-1})}_{e_t} = \ell_{t-1} + \alpha e_t$$

where  $e_t = y_t - \ell_{t-1}$  is the one-step-ahead forecast error. (SES is equivalent to ARIMA(0,1,1) — see Lecture 4.)

At each period the level is **updated in proportion to the latest error**.  $\alpha$  controls the speed of adaptation.

### Large $\alpha$ (e.g. 0.8):

- Reacts quickly to shocks
- Noisy, volatile forecasts
- Good for rapidly changing series

### Small $\alpha$ (e.g. 0.1):

- Slow adaptation
- Smooth, stable forecasts
- Good for slowly evolving, near-stationary series

---

**Parameter estimation:** choose  $\alpha$  and  $\ell_0$  to minimize the sum of squared one-step-ahead errors (Hyndman and Athanasopoulos 2021, Ch. 8):

$$\min_{\alpha, \ell_0} \sum_{t=1}^T (y_t - \ell_{t-1})^2, \quad 0 < \alpha \leq 1$$

Solved numerically (bounded nonlinear optimization).

**Forecast:** for any horizon  $h \geq 1$ ,

$$\hat{y}_{T+h|T} = \ell_T$$

SES produces a **flat forecast** — the same value at all horizons. Appropriate only when the series has *no trend and no seasonality*.

*Socratic:* if  $y_T = 110$ ,  $\ell_{T-1} = 104$ , and  $\hat{\alpha} = 0.6$ , what is  $\hat{y}_{T+1|T}$ ?

Estimated  $\hat{\alpha} \approx 0.70$  — fast adaptation. Forecast for month  $T+1$  equals the final smoothed level  $\ell_T$ . Lab 3 compares RMSE against Lecture 1–2 benchmarks.

### Connection to benchmarks:

Special case	Model
$\alpha = 1$	Naïve forecast: $\hat{y}_{T+h T} = y_T$
$\alpha \rightarrow 0$ (large $T$ )	Historical mean: $\hat{y}_{T+h T} = \bar{y}$

SES **nests** both extreme benchmarks from Lecture 1.  $\hat{\alpha}$  locates the optimal point between “forget everything” and “forget nothing.”

---

## Holt's Linear Method

Add a second component for trend: level and slope updated separately.

---

**Level:**  $\ell_t = \alpha y_t + (1 - \alpha)(\ell_{t-1} + b_{t-1})$

**Trend:**  $b_t = \beta^*(\ell_t - \ell_{t-1}) + (1 - \beta^*)b_{t-1}$

Parameters:  $0 < \alpha \leq 1$ ,  $0 < \beta^* \leq 1$  (Holt 2004)

### Interpretation:

- $\ell_t$ : where the series is *now*
- $b_t$ : the current growth rate
- $\beta^*$  controls how fast the slope adapts  
*SES is nested inside Holt: take  $\beta^* \rightarrow 0$  and  $b_0 = 0$  (so  $b_t = 0$  for all  $t$ ).*

### $h$ -step forecast:

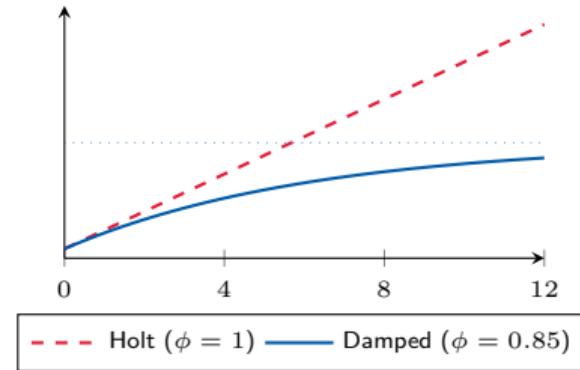
$$\hat{y}_{T+h|T} = \ell_T + h b_T$$

Extrapolates the current trend linearly.

Linear extrapolation can be **unrealistic** for long horizons — trends rarely persist indefinitely.

**Damped-trend modification (Gardner 2006):**

$$\hat{y}_{T+h|T} = \ell_T + \sum_{j=1}^h \phi^j b_T, \quad 0 < \phi < 1$$



$\phi \rightarrow 1$ : standard Holt (no damping).

$\phi \rightarrow 0$ : forecast  $\rightarrow \ell_T$  (flat; trend vanishes).

**Use damped by default for  $h > 6$ .**

*Socratic:* with  $\ell_T = 100$ ,  $b_T = 4$ ,  $\phi = 0.85$ , what is  $\hat{y}_{T+3|T}$ ? (Hint: compute  $\phi + \phi^2 + \phi^3$ .) Why might large  $\beta^*$  cause erratic forecasts when spikes occur?

---

## Holt-Winters Seasonal Method

Add a third component for seasonality: level, trend, and season.

---

---

The **additive Holt-Winters model** (Winters 1960):

$$y_t = \ell_t + b_t + s_t + \varepsilon_t$$

where  $s_t$  is the seasonal component with period  $m$ .

$$\ell_t = \alpha(y_t - s_{t-m}) + (1 - \alpha)(\ell_{t-1} + b_{t-1})$$

$$b_t = \beta^*(\ell_t - \ell_{t-1}) + (1 - \beta^*) b_{t-1}$$

$$s_t = \gamma(y_t - \ell_{t-1} - b_{t-1}) + (1 - \gamma) s_{t-m}$$

Parameters:  $0 < \alpha \leq 1$ ,  $0 < \beta^* \leq 1$ ,  $0 < \gamma \leq 1 - \alpha$ ; seasonal period  $m$

*The  $m$  seasonal indices sum to zero over one full period. For monthly data,  $m = 12$ .*

---

### Additive forecast:

$$\hat{y}_{T+h|T} = \ell_T + h b_T + s_{T+h-m}$$

### Multiplicative forecast:

$$\hat{y}_{T+h|T} = (\ell_T + h b_T) s_{T+h-m}$$

Variant	Seasonal component	When to use
Additive	$\ell_t + s_t$	Seasonal amplitude <i>constant</i> in size
Multiplicative	$\ell_t \times s_t$	Seasonal amplitude <i>proportional</i> to level

The holiday-season spike grows as overall sales grow  $\Rightarrow$  **multiplicative seasonality** is usually preferred.

---

**Parameters to estimate:**  $\alpha$ ,  $\beta^*$ ,  $\gamma$ , plus initial values  $\ell_0$ ,  $b_0$ , and  $m$  seasonal indices  $s_{1-m}, \dots, s_0$ .

### Initialization strategies:

- **Classical:** first-season average for  $\ell_0$ ; first-to-second-season slope for  $b_0$ ; seasonal differences for  $s_1, \dots, s_m$
- **Optimized:** jointly minimize SSE over all parameters *and* initial values (default in statsmodels)

Holt-Winters requires at least **2m** observations to initialize. With  $m = 12$ , you need  $\geq 24$  months of data.

*Socratic: if the December spike triples over ten years, which variant is appropriate and what does that imply for  $\hat{\gamma}$ ?*

---

## The ETS Framework

A unified state space taxonomy: Error  $\times$  Trend  $\times$  Seasonal.

---

---

Hyndman and Athanasopoulos (2021) unify all exponential smoothing variants as **state space models** with a single error source.

Component	Options	Meaning
Error ( $E$ )	A, M	Additive / Multiplicative
Trend ( $T$ )	N, A, $A_d$	None / Additive / Additive damped
Seasonal ( $S$ )	N, A, M	None / Additive / Multiplicative

$2 \times 3 \times 3 = 18$  combinations; some are inadmissible, giving **15 valid ETS models**.

**ETS(A,N,N)** = SES

**ETS(A,A,N)** = Holt linear

**ETS(A,A<sub>d</sub>,N)** = Holt damped

**ETS(A,A,A)** = Holt-Winters additive

**ETS(A,A,M)** = Holt-Winters multiplicative  
(and 10 more variants)

Every ETS model can be written as a **state space model** (additive-error form shown):

$$y_t = h(\mathbf{x}_{t-1}) + k(\mathbf{x}_{t-1}) \varepsilon_t, \quad \varepsilon_t \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \sigma^2)$$

$$\mathbf{x}_t = f(\mathbf{x}_{t-1}) + g(\mathbf{x}_{t-1}) \varepsilon_t$$

where  $\mathbf{x}_t = (\ell_t, b_t, s_t, \dots)'$  is the state vector.

#### Additive error ( $E = A$ ):

- Gaussian likelihood, closed-form prediction interval (PI)
- Easier to interpret

#### Multiplicative error ( $E = M$ ):

- Often better RMSE
- No closed-form PI

The state space form enables **maximum likelihood estimation** and principled prediction intervals.

Notation:  $\varepsilon_t$  here is the model innovation. For additive-error models,  $\varepsilon_t \equiv e_t$  (forecast error); for multiplicative-error models,  $\varepsilon_t = e_t / \ell_{t-1}$ .

Fit all 15 valid ETS models; select by minimum AIC. (Hyndman and Athanasopoulos 2021, Ch. 8):

$$AIC = -2\hat{L} + 2p, \quad \hat{L} = \text{maximized log-likelihood}$$

ETS model	Key features	Typical application
ETS(A,N,N)	No trend, no seasonal	Stable, level series
ETS(A,A <sub>d</sub> ,N)	Damped trend	Long-horizon trending
ETS(A,A,M)	Linear trend, mult. seasonal	Retail, tourism

In the M4 Competition, ETS-family methods ranked among the most accurate purely statistical approaches (Makridakis et al. 2020).

Lab 3 uses `statsmodels.tsa.exponential_smoothing.ets.ETSMModel` for full AIC-based model selection across all valid ETS variants.

---

## Key Takeaways and Roadmap

ETS unifies exponential smoothing as state space models with AIC selection.

---

---

**Exponential weighting** gives more importance to recent observations;  $\alpha$  controls the forgetting rate.

**SES**: one smoothing parameter ( $\alpha$ ), flat forecast — nests naïve ( $\alpha=1$ ) and mean ( $\alpha\rightarrow 0$ ).

**Holt**: adds trend component  $b_t$  — use the *damped* variant for horizons  $h > 6$ .

**Holt-Winters**: adds seasonal component  $s_t$  — choose additive vs. multiplicative based on whether amplitude scales with the level.

**ETS framework**: unifies all methods as state space models; AIC selects the best variant automatically.

*All ETS models share the error-correction intuition: update each component in proportion to the latest forecast error.*

---

Exponential smoothing captures **level, trend, and seasonality** through adaptive updating.

### What ETS cannot easily handle:

- Rich autocorrelation structure beyond a single AR(1)-like decay
- Integration and differencing in a principled statistical framework
- Combining external regressors with adaptive dynamics (ARIMAX)

**Lecture 4:** Stationarity, differencing, and the ARIMA model family — a regression-based approach to capturing autocorrelation.

**Lab 3:** Fit SES, Holt, Holt-Winters, and auto-ETS to RSXFS; compare RMSE / MAE with Lecture 1–2 benchmarks.

- 
-  Gardner, Everette S. (2006). "Exponential Smoothing: The State of the Art—Part II". In: *International Journal of Forecasting* 22.4, pp. 637–666.
  -  Holt, Charles C. (2004). "Forecasting Seasonals and Trends by Exponentially Weighted Moving Averages". In: *International Journal of Forecasting* 20.1. Reprint of the 1957 ONR Research Memorandum, pp. 5–10.
  -  Hyndman, Rob J. and George Athanasopoulos (2021). *Forecasting: Principles and Practice*. 3rd. Open access. Melbourne, Australia: OTexts. URL: <https://otexts.com/fpp3/>.
  -  Makridakis, Spyros, Evangelos Spiliotis, and Vassilios Assimakopoulos (2020). "The M4 Competition: 100,000 Time Series and 61 Forecasting Methods". In: *International Journal of Forecasting* 36.1, pp. 54–74.
  -  Winters, Peter R. (1960). "Forecasting Sales by Exponentially Weighted Moving Averages". In: *Management Science* 6.3, pp. 324–342.