

---

# Machine Learning Introduction

## BSAD 8310: Business Forecasting — Lecture 7

---

Department of Economics  
University of Nebraska at Omaha  
Spring 2026

- 
- 1 Why ML for Forecasting?**
  - 2 The Bias-Variance Tradeoff**
  - 3 Train / Validation / Test Discipline**
  - 4 Cross-Validation for Time Series**
  - 5 Regularization Preview**
  - 6 Feature Engineering Preview**
  - 7 The ML Forecasting Pipeline**

---

## Why ML for Forecasting?

Classical methods are powerful — but structured. What happens when the structure breaks?

---

---

Six lectures built a powerful toolkit. Now we stress-test its assumptions.

## What we have (L01–L06):

- Benchmarks and regression (L01–L02)
- ETS for trend + seasonality (L03)
- ARIMA/SARIMA for autocorrelation (L04)
- VAR, ARIMAX, ECM for multivariate dynamics (L05)
- Rigorous walk-forward evaluation (L06)

Classical models assume: linearity, fixed lag structure, Gaussian errors, known seasonality. When these fail, **model misspecification** inflates forecast error beyond what better data could fix.

*Socratic: which assumption fails first when forecasting social-media-driven demand spikes for a new product?*

Gap	Classical limit	ML remedy
Non-linearity	Linear $f(x)$ assumed	Trees, nets
Many predictors	VAR fails ( $p > T$ )	Regularization
Interactions	Hand-specified only	Learned
Non-Gaussian	Normality assumed	Distribution-free

ML does **not** replace ARIMA. It extends the toolkit for problems where: (1) predictors number in the hundreds, (2) relationships are non-linear, or (3) structure is unknown *a priori*.

(Hastie et al. 2009): the statistical learning perspective on non-linearity and interactions.

---

**100,000 time series, 61 methods** (Makridakis et al. 2020): nine of the top twelve methods were *hybrid* (classical + ML).

Method	Rank	sMAPE
Hybrid ES-RNN (ML + ETS)	1	11.37
Theta (classical)	2	11.37
FFORMA (ML ensemble)	3	11.57
Naive 2 (seasonal)	12	13.56
Auto ARIMA	14	13.58

Retail, electricity, and supply-chain forecasting now routinely use gradient-boosted trees and ensemble combinations of ARIMA + ML.

But: pure ML often *underperforms* naive on short series ( $n < 100$ ). Context matters.

---

## The Bias-Variance Tradeoff

Every modelling decision shifts error between bias and variance. This is the central tension in all of statistical learning.

---

Let  $\hat{f}$  be an estimated model and  $f$  the true function. At a new point  $x_0$ :

$$\text{MSE}(\hat{f}(x_0)) = (\text{Bias}[\hat{f}(x_0)])^2 + \text{Var}[\hat{f}(x_0)] + \sigma^2$$

$\sigma^2$  = irreducible noise (no model eliminates this).

- **Bias**<sup>2</sup>: systematic error from wrong model class (underfitting)
- **Variance**: sensitivity to particular training sample (overfitting)
- $\sigma^2$ : irreducible — better models cannot reduce this, only better data can

$\hat{f}$ : estimated model.  $f$ : true unknown function.  $x_0$ : feature vector at new observation.  $\sigma^2 = \text{Var}[\varepsilon]$  where  $\varepsilon$  is the noise term.

## Both extremes hurt.

AR(1) fit to a threshold non-linear series: the model systematically underestimates peaks and overestimates troughs every cycle.

Low variance: refit on any subsample and the errors look the same.

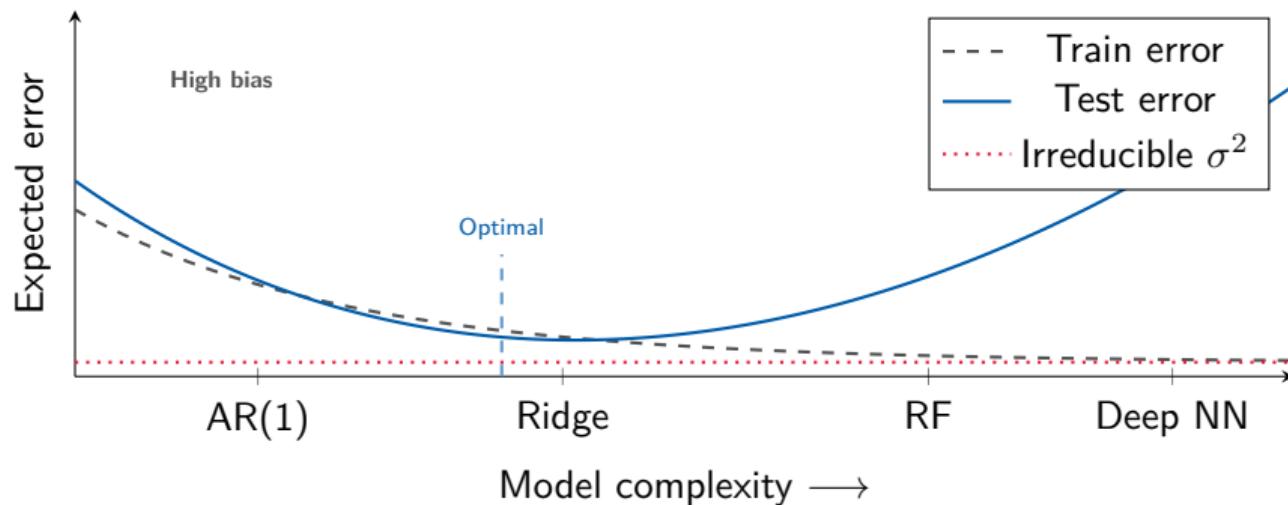
High variance — overfitting:

Degree-15 polynomial fits every wiggle in the training set. Train RMSE  $\approx 0.02$ . Test RMSE  $\approx 8 \times$  train RMSE.

Small change in training data  $\Rightarrow$  large change in  $\hat{f}$ .

Reducing bias *typically* increases variance. We control the tradeoff via **model complexity** and **regularization**.

Train error always falls with complexity. Test error has a U-shape.



*Cross-validation finds the optimal complexity point without ever touching the test set.*

Model	Bias	Var.	Good when
AR(1)	High	Low	Long, stable series
ARIMA (auto)	Medium	Low–Med	General TS
Ridge (many X)	Med	Low	Many predictors
OLS (unregul.)	Low	High	$n \gg p$ only
Deep LSTM	Low	High	Very long series

*This directly motivates LASSO (L08): deliberately increasing bias via shrinkage to reduce variance on short, noisy series.*

For short series ( $n < 200$ ):

**Bias is cheaper than variance.**

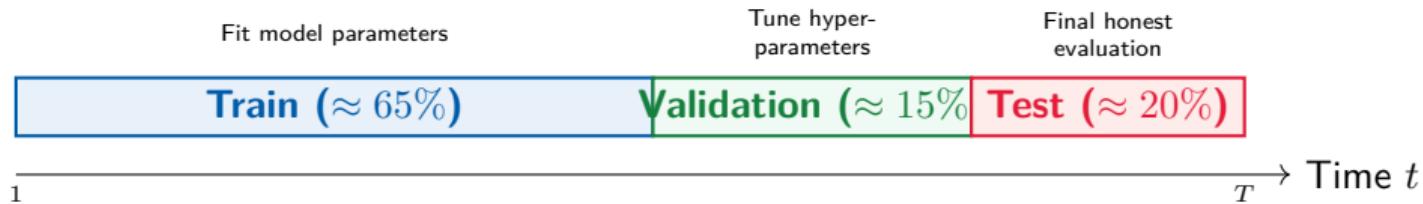
Prefer parsimonious models. Add complexity only when cross-validation confirms out-of-sample improvement.

---

## Train / Validation / Test Discipline

One split is not enough. Three sets, three roles — and the rules are strict.

---



## Why three sets, not two?

If you tune on the test set — even once — it becomes a second validation set, and you no longer have an honest generalization estimate.

The test set is a **time capsule**. Open it exactly once, at the very end, to report the final number.

Leakage occurs when information from outside the training window contaminates the model, producing optimistic in-sample performance that does not generalize out-of-sample.

### Feature leakage:

- Using next-month CPI to forecast this-month sales (future predictor)
- Target-encoding computed on the full dataset before splitting (statistical leakage)

Random train/test split is *correct* for i.i.d. data. It is **wrong** for time series. Always split in **chronological order**.

In production: leakage makes reported RMSE optimistic by  $2\times\text{--}5\times$ .

### Temporal leakage:

- Random train/test shuffle on a time series
- Fitting StandardScaler on the full series before splitting

1. **No random shuffling** — always chronological order
2. **Validation after train, test after validation** — strict temporal ordering; no overlaps
3. **No future-dependent features** — any feature using time  $t + k$  data cannot be known at time  $t$
4. **Scale within each fold** — fit StandardScaler on train-fold only; transform val/test with train parameters
5. **Optional gap** — buffer between train end and val start prevents autocorrelation bleed-through

These rules are the time-series extensions of the bias-variance discipline.

Break them and every accuracy number you report is invalid.

`train_test_split(shuffle=False)` respects ordering. Use

TimeSeriesSplit for CV.

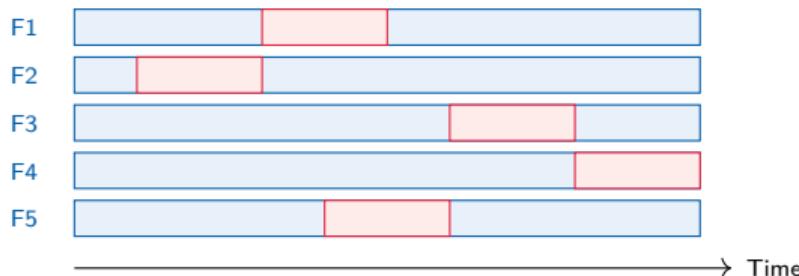
---

## Cross-Validation for Time Series

Standard  $k$ -fold CV violates time ordering. Walk-forward CV is the fix.

---

## Standard 5-fold (shuffled):



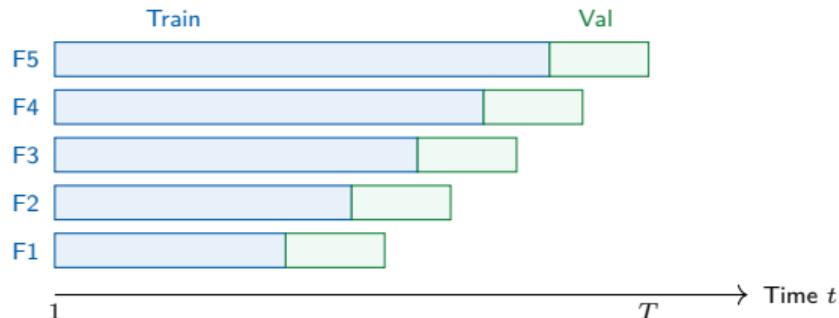
Wrong for time series: validation block may precede some training data  $\Rightarrow$  temporal leakage.

## Failure modes:

- Uses future data to predict the past
- Train and validation observations are not exchangeable under autocorrelation
- Validation RMSE is optimistically biased

Bergmeir et al. (2018): random  $k$ -fold gives biased CV estimates for autoregressive series. The bias grows with autocorrelation strength.

**The correct analog:** expanding-window validation introduced in L06, now applied to hyperparameter tuning.



```
from sklearn.model_selection import TimeSeriesSplit  
tscv = TimeSeriesSplit(n_splits=5, gap=0)  
for train_idx, val_idx in tscv.split(X):  
    X_train, X_val = X[train_idx], X[val_idx]
```

gap: skip observations between train end and val start. Set gap= $H$  when forecasting  $H$  steps ahead to prevent overlap.

General CV theory: (Arlot and Celisse 2010).

- 
1. Choose candidate values:  
 $\lambda \in \{0.01, 0.1, 1, 10, 100\}$
  2. For each  $\lambda$ : run TimeSeriesSplit CV,  
compute mean validation RMSE
  3. Select  $\lambda^*$  minimising mean CV RMSE
  4. **Refit on full train + val** with  $\lambda^*$
  5. Evaluate exactly once on held-out test set

**Step 4 is critical:** refit on all non-test data.  
Using the CV train-fold only wastes data.

5-fold TimeSeriesSplit:  
 $\lambda = 0.1 \Rightarrow \text{CV RMSE} = 1,780$  (optimal)  
 $\lambda = 0.001$  (overfit): 1,890  
 $\lambda = 100$  (underfit): 2,140  
Test RMSE with  $\lambda^* = 0.1$ : 1,620.

---

## Regularization Preview

When there are many predictors, unconstrained OLS overfits. Shrinkage is the solution. Full treatment: Lecture 8.

---

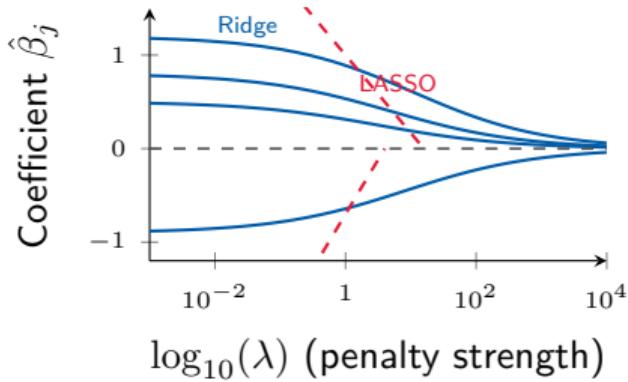
2009

$$\hat{\beta} = \arg \min_{\beta} \underbrace{\sum_{t=1}^T (y_t - \beta^\top x_t)^2}_{\text{OLS loss}} + \lambda \cdot P(\beta)$$

**Ridge:**  $P(\beta) = \|\beta\|_2^2$     **LASSO (Tibshirani 1996):**  $P(\beta) = \|\beta\|_1$     **EN (Zou and Hastie 2005):**  $\alpha\|\beta\|_1 + (1-\alpha)\|\beta\|_2^2$

- **Ridge:** shrinks all coefficients smoothly; none reach exactly zero
- **LASSO:** sets some coefficients to exactly zero  $\Rightarrow$  automatic variable selection
- **Elastic Net:** handles collinear predictors; ridge + LASSO combined

All three trade bias for variance via  $\lambda$ .  
 $\lambda = 0$ : OLS (maximum variance).  
 $\lambda \rightarrow \infty$ : intercept only (maximum bias).  
Full derivation in L08.



As  $\lambda \rightarrow \infty$ :  $\hat{\beta}_j \rightarrow 0$  (all predictors removed).

As  $\lambda \rightarrow 0$ : OLS solution.

LASSO paths hit zero at a finite  $\lambda$  — Ridge paths only approach zero asymptotically.

CV selects  $\lambda^*$  at the bias-variance minimum.

*Illustrative paths. L08 derives the exact form from the subgradient conditions.*

---

## Feature Engineering Preview

ML models do not handle time series natively. We must create temporal structure explicitly.  
Full treatment: Lecture 11.

---

A raw series  $y_1, \dots, y_T$  is not a feature matrix. We extract temporal patterns as columns.

Feature type	Formula	Business meaning
Lag	$x_t^{(k)} = y_{t-k}$	Sales $k$ periods ago
Rolling mean	$\bar{y}_{t,w} = \frac{1}{w} \sum_{k=1}^w y_{t-k}$	Recent trend level
Rolling std	$s_{t,w} = \text{std}(y_{t-w:t-1})$	Recent volatility
Month-of-year	$\mathbf{1}[\text{month}(t) = m]$	Seasonal dummy
Trend counter	$t = 1, 2, \dots, T$	Linear time trend

All features **must** use only  $y_1, \dots, y_{t-1}$  when predicting  $y_t$ . Any feature using  $y_t$  or later creates leakage. Apply `.shift(1)` before any rolling window.

*Full treatment in L11: calendar features, Fourier terms, interaction features, and pipeline automation.*

```
# Lag features: shift(1) avoids leakage
X['lag_1'] = y.shift(1)
X['lag_12'] = y.shift(12)

# Rolling mean of  $y_{t-1} \dots y_{t-3}$ 
X['roll_3'] = y.shift(1).rolling(3).mean()
()

# Calendar feature (always known)
X['month'] = y.index.month

# Drop rows with NaN from lags
X = X.dropna()
```

- `.shift(1)`: shifts series forward one period — ensures  $x_t$  uses  $y_{t-1}$ , not  $y_t$
- `.rolling(3).mean()` after `.shift(1)`: mean of  $y_{t-1}, y_{t-2}, y_{t-3}$
- `.index.month`: calendar feature; no leakage (calendar is always known in advance)

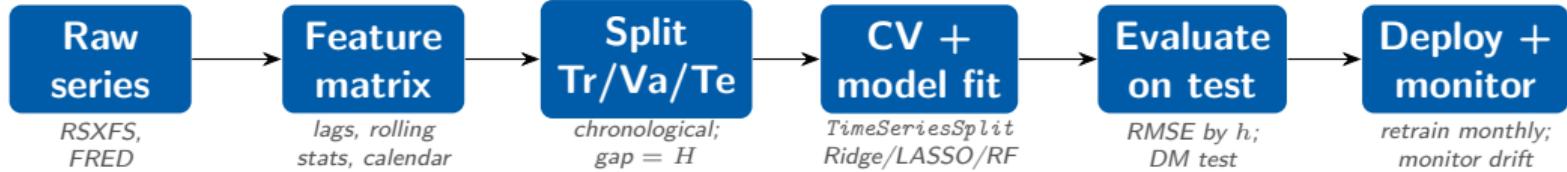
L11 wraps this logic in a `sklearn.Pipeline` that prevents leakage automatically at every CV fold.

---

## The ML Forecasting Pipeline

End-to-end: features → split → fit → evaluate → deploy.

---



The `sklearn.Pipeline` object chains preprocessing (`StandardScaler`) + model into a single estimator. This ensures the scaler is fitted on train-fold only during CV — eliminating a common source of leakage (James et al. 2021).

ML **extends** classical forecasting for non-linearity, many predictors, and unknown structure.

$\text{MSE} = \text{Bias}^2 + \text{Var} + \sigma^2$  governs every modelling decision.

Three-way train/val/test split with chronological ordering is non-negotiable.

Walk-forward CV (`TimeSeriesSplit`) replaces  $k$ -fold for time series.

Regularization (L08) and feature engineering (L11) are the main levers for controlling bias-variance.

## Lecture roadmap:

Lecture	Topic	Lab 7:
L08	LASSO, Ridge, EN	
L09	RF, XGBoost	
L10	LSTM, attention	
L11	Feature pipeline	
L12	Capstone cases	

*full pipeline from raw RSXFS to LASSO and Ridge with walk-forward CV.*

- 
-  Arlot, Sylvain and Alain Celisse (2010). "A Survey of Cross-Validation Procedures for Model Selection". In: *Statistics Surveys* 4, pp. 40–79.
  -  Bergmeir, Christoph, Rob J. Hyndman, and Bonsoo Koo (2018). "A Note on the Validity of Cross-Validation for Evaluating Autoregressive Time Series Prediction". In: *Computational Statistics & Data Analysis* 120, pp. 70–83.
  -  Hastie, Trevor, Robert Tibshirani, and Jerome Friedman (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. 2nd. New York: Springer.  
URL: <https://hastie.su.domains/ElemStatLearn/>.
  -  James, Gareth et al. (2021). *An Introduction to Statistical Learning with Applications in Python*. 2nd. New York: Springer. URL: <https://www.statlearning.com/>.
  -  Makridakis, Spyros, Evangelos Spiliotis, and Vassilios Assimakopoulos (2020). "The M4 Competition: 100,000 Time Series and 61 Forecasting Methods". In: *International Journal of Forecasting* 36.1, pp. 54–74.

- 
-  Tibshirani, Robert (1996). "Regression Shrinkage and Selection via the Lasso". In: *Journal of the Royal Statistical Society: Series B* 58.1, pp. 267–288.
  -  Zou, Hui and Trevor Hastie (2005). "Regularization and Variable Selection via the Elastic Net". In: *Journal of the Royal Statistical Society: Series B* 67.2, pp. 301–320.