

---

# Lecture 09: Tree-Based Methods

## Random Forests and Gradient Boosting for Forecasting

---

BSAD 8310: Business Forecasting

University of Nebraska at Omaha

Spring 2026

- 
- 1 Decision Trees: CART**
  - 2 Bagging and Random Forests**
  - 3 Gradient Boosting and XGBoost**
  - 4 Feature Importance**
  - 5 Hyperparameter Tuning**
  - 6 Application to Forecasting**
  - 7 Takeaways and References**

---

**Lecture 08 best result:** Elastic Net, RMSE  $\approx$  2,540 on RSXFS.

Elastic Net is a linear model:  $\hat{y} = \mathbf{x}^\top \hat{\boldsymbol{\beta}}$ . It cannot capture *interactions* or *threshold effects* (e.g., the holiday effect amplifies when the economy is strong).

### What tree-based methods add:

- **Nonlinearity:** predictions are piecewise constants over feature regions
- **Interactions:** splits on  $x_1$  in a region defined by  $x_2$  capture  $x_1 \times x_2$  interactions automatically
- **No standardization required:** scale-invariant splits

**Today's goal:** build from a single decision tree to Random Forests and XGBoost, then compare against the Lecture 08 leaderboard.

---

**Decision trees** partition the feature space into rectangular regions using recursive binary splits. They are the building blocks of Random Forests and gradient boosting — understanding them is essential before ensembles.

**CART (Classification and Regression Trees) for regression:** at each node, find the split  $(j, t)$  minimizing the weighted within-region variance:


$$\min_{j, t} \left[ \sum_{i: x_{ij} \leq t} (y_i - \bar{y}_L)^2 + \sum_{i: x_{ij} > t} (y_i - \bar{y}_R)^2 \right]$$

### Algorithm:

1. Start with all data at the root
2. For every feature  $j$  and every threshold  $t$ :  
compute the split criterion
3. Split on the  $(j^*, t^*)$  giving the largest reduction
4. Recurse on left and right children
5. Stop when nodes are too small or depth limit reached
6. Predict:  $\hat{y} = \bar{y}_{\text{leaf}}$

Each split depends on which points fall into a region. Optimal splitting is NP-hard globally; greedy top-down splitting is the standard approximation (CART).

## Example: 3-leaf tree for RSXFS



## Key properties:

- **Non-parametric** — no functional form assumed
- **Nonlinear** — can capture threshold and interaction effects
- **Interpretable** — decision path is readable
- **Handles mixed features** — lags, rolling stats, dummies all treated equally
- **Scale-invariant** — no need to standardize predictors


A single deep tree **overfits severely**. The solution is not pruning alone — it is *ensembling*.

## Tree depth controls complexity:

- **Depth 1** (stump): high bias, low variance
- **Depth 5–10**: balanced region
- **Full tree** (each leaf = 1 obs): zero bias, infinite variance

## Why trees have high variance:

A single split can change the entire subtree beneath it. Small perturbations in data (e.g., one outlier) produce very different trees — trees are *unstable*.



**Key insight:** averaging many high-variance, low-bias trees (each fit to a bootstrap sample) substantially reduces variance while preserving low bias. This is bagging.

---

**Random Forests** reduce the variance of decision trees by: (1) training each tree on a different *bootstrap sample* (bagging), and (2) using a random *feature subset* at each split (decorrelation). The result is a low-variance, competitive forecaster (Breiman 2001).

1996

Draw  $B$  bootstrap samples; fit one deep tree on each; average:  $\hat{y}_{\text{bag}} = \frac{1}{B} \sum_{b=1}^B \hat{f}^{(b)}(\mathbf{x})$

### Variance of the ensemble:

$$\text{Var}(\bar{f}) = \rho\sigma^2 + \frac{1 - \rho}{B}\sigma^2$$

$\rho$  = pairwise prediction correlation between trees;  $\sigma^2$  = single-tree variance.

As  $B \rightarrow \infty$ , variance  $\rightarrow \rho\sigma^2$  (not zero). Reducing  $\rho$  matters as much as increasing  $B$  — this motivates Random Forests.

### Bagging key numbers:

Bootstrap sample  $\approx 63\%$  unique obs.

Remaining  $\approx 37\% =$  out-of-bag (OOB)

OOB obs. give a free cross-validation estimate

$B \geq 500$  trees recommended (stable)

**Problem with bagging:** if one feature dominates, all trees use it at the root  $\Rightarrow$  trees are highly correlated ( $\rho$  stays large).

**Random Forest fix:** at each split, randomly sample  $m < p$  features and find the best split *among those  $m$  only*.

- Default:  $m = \lfloor \sqrt{p} \rfloor$  (sklearn  $\geq 1.1$ );  $\lfloor p/3 \rfloor$  recommended for regression (Hastie et al.)
- Reduces  $\rho \Rightarrow$  lower ensemble variance
- Each tree is weaker but the ensemble is stronger

RF = Bagging + random feature selection at each split. Both ingredients are necessary.

Each obs. is OOB for  $\approx 37\%$  of trees.  
Predict using only those trees; compute RMSE.

OOB RMSE  $\approx$  5-fold CV RMSE at large  $B$  — essentially *free* cross-validation.

---

Parameter	sklearn name	Default	Effect
Num. trees	<code>n_estimators</code>	100	More = lower variance; plateau after $\sim 500$
Max features/split	<code>max_features</code>	'sqrt'	Lower = more decorrelated; tune via OOB
Max depth	<code>max_depth</code>	None	None = full tree; limit to reduce memory
Min samples/leaf	<code>min_samples_leaf</code>	1	Higher = smoother predictions; reduces overfit
Bootstrap	<code>bootstrap</code>	True	False = subsampling (not bootstrap)

*Socratic: why does increasing `n_estimators` always reduce out-of-bag error, but increasing tree depth does not? (Hint: think about what each controls in the bias-variance decomposition.)*

---

**Gradient boosting** builds trees *sequentially*: each new tree fits the residuals left by the previous ensemble. XGBoost adds regularization and second-order optimization to this framework (Chen and Guestrin 2016).

Initialize  $F_0(\mathbf{x}) = \bar{y}$ . For  $b = 1, \dots, B$ :

- Residuals:  $r_i^{(b)} = y_i - F_{b-1}(\mathbf{x}_i)$  (negative gradient for squared-error loss)
- Fit shallow tree  $T_b$  to  $\{(\mathbf{x}_i, r_i^{(b)})\}$
- Update:  $F_b = F_{b-1} + \eta T_b$

Predict:  $\hat{y} = F_B(\mathbf{x})$

## RF vs. Boosting:

- RF: **parallel** trees, deep, bias not reduced by adding more trees
- Boosting: **sequential**, shallow trees (depth 3–6)
- Boosting bias **decreases** with stages
- Both benefit from large  $B$  (number of trees)

Small  $\eta$  + large  $B$  + early stopping outperforms large  $\eta$  + small  $B$ .

**Always use early stopping on a validation set.**

---

XGBoost (Chen and Guestrin 2016) extends gradient boosting with three additions:

**1. Newton step (2nd-order):**

$$\mathcal{L}^{(b)} = \sum_i [g_i f_b(\mathbf{x}_i) + \frac{1}{2} h_i f_b^2(\mathbf{x}_i)] + \Omega(f_b)$$

$g_i, h_i$  = gradient and Hessian of loss.

**2. Explicit regularization:**

$$\Omega(f_b) = \gamma T + \frac{1}{2} \lambda \|\mathbf{w}\|_2^2$$

$T$  = leaf count,  $\mathbf{w}$  = leaf weights,  $\gamma$  = per-leaf complexity penalty.

**3. Column subsampling:**

Like RF, randomly sample features per tree *and* per split — reduces correlation and overfitting.

Newton step  $\Rightarrow$  more precise gradient direction. Regularization  $\Rightarrow$  guards against overfitting. Together: XGBoost beats vanilla GBM on most benchmarks.

*Note:  $\lambda$  here is XGBoost's L2 leaf-weight penalty (distinct from the regularization  $\lambda$  in Lecture 08).*

Parameter	XGB name	Typical	Effect
Num. trees	n_estimators*	500–2000	More + small $\eta$ = better (use early stopping)
Learning rate	learning_rate	0.01–0.1	Smaller = more robust; requires more trees
Tree depth	max_depth	3–6	Shallow preferred; deeper = overfit risk
Subsample	subsample	0.7–0.9	Row fraction per tree; adds stochasticity
Col. sample	colsample_bytree	0.7–0.9	Feature fraction per tree (like RF)
L2 penalty	reg_lambda	1	Ridge on leaf weights; stabilizes

\*sklearn API name; native API uses num\_boost\_round.

**Practical starting point:** learning\_rate=0.05, max\_depth=4, n\_estimators=1000 with early stopping on val RMSE. Tune max\_depth and subsample via TimeSeriesSplit CV.

---

Tree-based models offer built-in feature importance measures. Understanding *which features drive predictions* is essential for business interpretation and debugging.

$$\text{Imp}(j) = \sum_{\text{nodes using } j} \Delta \text{RSS} \cdot \frac{n_{\text{node}}}{n}$$

Sum of RSS reduction at all splits on feature  $j$ , weighted by fraction of samples reaching that node.

**Caution:** biased toward *high-cardinality* features (continuous variables get more split opportunities than binary dummies). Do not use for final reporting.

For each feature  $j$ :

1. Shuffle column  $j$  in the OOB set
2. Record increase in RMSE

Larger increase  $\Rightarrow$  more important.

**Unbiased** (no cardinality bias).

**Model-agnostic.** Use for reporting.

Average *gain* (improvement in objective function) per split using feature  $j$ , across all trees.

More interpretable than impurity for boosted trees because it weights by actual objective reduction.

### Practical rule:

- Use **permutation** importance for model reports and presentations
- Use **impurity** for quick diagnostics during training
- Use **XGBoost gain** when comparing feature contributions inside a boosted model

1. **lag\_12** — dominant seasonal signal
2. **lag\_1** — short-run momentum
3. **roll\_mean\_12** — trend level
4. **lag\_3** — quarterly effect
5. **month\_12** — December holiday

Consistent across all three importance types — feature engineers should lead with **lag\_12**.

---

For time series, hyperparameter tuning must use `TimeSeriesSplit` to respect temporal ordering. We can use `GridSearchCV` (exhaustive) or `RandomizedSearchCV` (efficient for large grids).

## Random Forest (RandomizedSearchCV):

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import (
    TimeSeriesSplit, RandomizedSearchCV)

tscv = TimeSeriesSplit(n_splits=5, gap=1) # gap>=1 for 1-step-
                                         ahead
rf_grid = {
    'n_estimators': [200, 500],
    'max_features': ['sqrt', 0.33],
    'min_samples_leaf': [1, 3, 5],
    'max_depth': [None, 10, 20],
}
rf = RandomizedSearchCV(
    RandomForestRegressor(random_state=42),
    rf_grid, n_iter=20,
    cv=tscv,
    scoring='neg_root_mean_squared_error',
    n_jobs=-1, random_state=42)
rf.fit(X_trainval, y_trainval)
```

## XGBoost (early stopping on val):

```
import xgboost as xgb

dtrain = xgb.DMatrix(X_train, label=y_train)
dval = xgb.DMatrix(X_val, label=y_val)
dtest = xgb.DMatrix(X_test, label=y_test)

params = {
    'learning_rate': 0.05,
    'max_depth': 4,
    'subsample': 0.8,
    'colsample_bytree': 0.8,
    'reg_lambda': 1.0,
    'objective': 'reg:squarederror',
}
model = xgb.train(
    params, dtrain,
    num_boost_round=2000,
    evals=[(dval, 'val')],
    early_stopping_rounds=50)
```

---

Apply Random Forest and XGBoost to the RSXFS feature matrix from Lecture 08. Compare against SARIMA and the best regularized model (Elastic Net) from Lecture 08.

**Same feature matrix as Lecture 08:** 12 lags + 3 rolling windows + 11 month dummies = 26 features (one dummy omitted vs. L08 for comparability; trees do not require it). **Key differences vs. regularized regression:**

### Advantages of trees:

- Capture nonlinearities and interactions (e.g., “holiday effect only when economy is strong”)
- No standardization needed
- Handle irrelevant features gracefully (unused features just don’t appear in splits)
- Built-in feature importance

### Cautions for time series:

- **No extrapolation:** trees predict  $\bar{y}_{\text{leaf}}$ , so they cannot predict beyond the training range
- **Trending series:** a naive tree trained on non-stationary data may perform poorly out of sample
- **Fix:** use first-differences or add rolling statistics as features to encode trend

**Stationarity matters for trees too.** If  $y_t$  trends upward beyond the training max, leaf means will systematically under-predict. Difference or detrend before applying tree-based forecasting.

---

## Test-set results on RSXFS (24-month horizon):

Model	RMSE	MAE
Seasonal Naïve	4 210	3 120
SARIMA(1,1,1)(1,1,1) <sub>12</sub>	2 840	2 100
Elastic Net ( $\lambda^*$ )	2 540	1 890
Random Forest	2 380	1 760
XGBoost (early stop)	2 250	1 650

Values are illustrative. Strategy: direct multi-step — one model trained per horizon  $h = 1, \dots, 24$ .

## Why do trees win here?

- Nonlinear interaction between lag\_12 and calendar dummies (holiday amplification)
- XGBoost sequential residual fitting corrects patterns that Elastic Net misses
- RF OOB error  $\approx$  CV error at  $B = 500$

**Caveat:** gains depend heavily on feature quality. With a good feature set, Elastic Net and XGBoost are often competitive. Adding more features benefits trees more.

**Decision trees** partition feature space greedily; they are interpretable but unstable (high variance).

**Random Forests** (bagging + random features) reduce variance dramatically while keeping bias low. OOB error is a free CV estimate.

**Gradient boosting** reduces bias sequentially by fitting residuals. XGBoost adds L2 regularization and Newton optimization.

**Feature importance** helps identify which lags, rolling stats, and calendar features drive forecasts. Use permutation importance for unbiased estimates.

**Trees cannot extrapolate** — ensure features encode trend and seasonality so predictions stay within the training distribution.

**RF vs. XGBoost:** prefer RF for speed and robustness with limited tuning; prefer XGBoost when maximum accuracy matters and you have time to tune `learning_rate` and `max_depth`.

**Preview of Lecture 10:** Neural Networks — LSTM and attention mechanisms can model long-range temporal dependencies that trees cannot.

- 
-  Breiman, Leo (1996). "Bagging Predictors". In: *Machine Learning* 24.2, pp. 123–140.
  -  — (2001). "Random Forests". In: *Machine Learning* 45.1, pp. 5–32.
  -  Chen, Tianqi and Carlos Guestrin (2016). "XGBoost: A Scalable Tree Boosting System". In: *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 785–794.