

---

# Lecture 11: Feature Engineering

## Building Better Inputs for Forecasting Models

---

BSAD 8310: Business Forecasting  
University of Nebraska at Omaha  
Spring 2026

- 
- 1 Lag Features**
  - 2 Rolling Statistics**
  - 3 Calendar and Structural Features**
  - 4 Interaction and Ratio Features**
  - 5 Feature Selection**
  - 6 Pipeline Design**
  - 7 Application to Forecasting**
  - 8 Takeaways and References**

---

**Lecture 10 best result:** LSTM with 36 features achieved RMSE  $\approx 1,920$  — down from XGBoost's 2,250 (26 features). The improvement came mostly from *better inputs*, not a fancier model.

### What feature engineering adds:

- **Lag features:** give the model explicit access to past values, guided by ACF and PACF
- **Rolling statistics:** capture local trends, volatility, and momentum across multiple time scales
- **Calendar effects:** encode seasonality without requiring seasonal differencing
- **Interaction / ratio features:** year-over-year and month-over-month change rates that tree models cannot discover on their own
- **Pipeline design:** prevent data leakage by fitting transformations inside each cross-validation fold

**Lecture goal:** Build `make_features_extended()` (36 features), select the best subset, and update the Lectures 01–11 leaderboard.

---

## Lag Features

Converting time-series forecasting into a supervised regression problem

---


---

**Lag features** are the most direct way to give a tree-based or linear model access to the past. They convert a time-series forecasting problem into a standard supervised regression problem with  $y_{t-k}$  as predictors.

### Design questions:

- Which lags are worth including? (ACF and PACF guidance)
- How many lags before diminishing returns?
- How to avoid look-ahead leakage when computing lags?

**Key rule:** always use `shift(k)` with  $k \geq 1$  so that at prediction time for period  $t$ , only values observed at  $t-1, t-2, \dots$  are used.



## Reading the charts:

- **ACF:** slow decay → persistent trend; spike at lag 12 → seasonal pattern
- **PACF:** significant at lag 1 only → pure AR(1); significant at lag 12 → include lag\_12

Include `lag_k` if the PACF at lag  $k$  exceeds  $\pm 1.96/\sqrt{n}$  (dashed band). For monthly RSXFS: include lags 1, 2, 12. Add lags 3–6 for robustness.

*Socractic: ACF at lag 12 is strong, but PACF at lag 12 is also strong. Why include lag\_12 even in an ARIMA model that handles seasonality through differencing?*

## Safe: always shift(k)

```
import pandas as pd

# SAFE: shift(k) looks back k steps
# At time t we only see y[t-k]
df['lag_1'] = df['y'].shift(1)
df['lag_2'] = df['y'].shift(2)
df['lag_12'] = df['y'].shift(12)

# Rolling mean: shift FIRST
df['roll_mean_3'] = (
    df['y'].shift(1)
    .rolling(3).mean())

# No future information used.
print("Lag features are safe.")
```

## Bug: missing shift(1)

```
# BUG: rolling without shift(1)
# includes y[t] = current target!
df['roll_mean_3'] = (
    df['y'].rolling(3).mean())

# At t=5, rolling uses:
# y[3], y[4], y[5] <- leakage!
# y[5] is what we are predicting.

# This inflates in-sample R^2
# but collapses on test data.

# Rule: always shift before
# any window operation.
```

Rolling without `shift(1)` is the #1 feature engineering bug. It causes dramatic overfitting that only appears on the test set.

---

## Rolling Statistics

Local trends, volatility, and momentum over time

---

---

**Rolling statistics** capture local trends, volatility, and momentum over multiple time windows. They are especially powerful for tree-based models, which cannot express the concept of a “moving average” from raw lag features alone.

### Feature families:

- **Rolling mean** ( $w = 3, 6, 12$ ): local trend at multiple scales
- **Rolling std** ( $w = 3, 6, 12$ ): local volatility — useful for identifying calm vs. volatile regimes
- **Rolling min/max** ( $w = 3, 6, 12$ ): range and extremes
- **Exponentially weighted mean (EWM)**: gives more weight to recent observations than a simple average

All rolling operations must be computed on `df['value'].shift(1)` to prevent leakage (see the Lag Leakage slide).

## Rolling window operations:

Feature	Formula	Captures
roll_mean_w	$\bar{y}_{t-w:t-1}$	Local trend
roll_std_w	$s_{t-w:t-1}$	Volatility
roll_min_w	$\min(y_{t-w}, \dots, y_{t-1})$	Trough
roll_max_w	$\max(y_{t-w}, \dots, y_{t-1})$	Peak

- $w = 3$ : captures quarter-long momentum
- $w = 6$ : captures half-year cycles
- $w = 12$ : captures full-year seasonality

Use all three; let the model select via regularization or feature importance.

## Exponentially weighted mean:

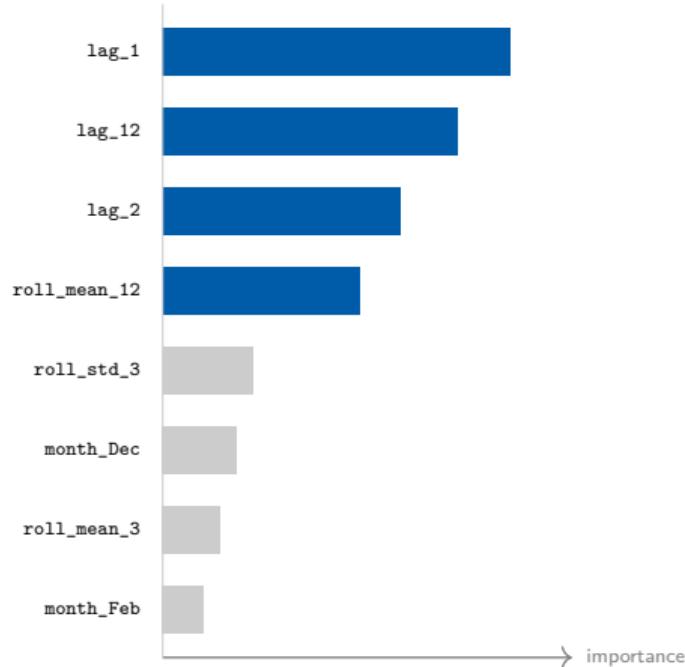
$$\text{EWM}_t = \alpha y_{t-1} + (1 - \alpha) \text{EWM}_{t-1}, \quad \alpha \in (0, 1)$$

With  $\alpha = 0.3$ : recent values receive  $\approx 3\times$  the weight of values 4 months back.

*Here  $\alpha$  is the EWM decay weight — distinct from the level-smoothing  $\alpha$  in Lecture 03 (ETS) and the L1/L2 mixing  $\alpha$  in Lecture 08 (Elastic Net).*

EWM reacts faster to trend reversals. On RSXFS, EWM with  $\alpha = 0.3$  reduces lag behind turning points by  $\sim 2$  months compared with a 12-month rolling mean.

## Permutation importance (Random Forest, val set):



### Key findings:

- Lags 1 and 12 dominate — confirming ACF/PACF
- `roll_mean_12` is the most important rolling feature
- `month_Dec` matters: December retail spike is real
- Calendar dummies are less important than rolling features when lags already capture seasonal levels

*This ranking is illustrative. Permutation importance is covered formally in Section 5 (Feature Selection).*

---

## Calendar & Structural Features

Encoding temporal position for tree-based models

---

---

**Calendar features** encode the position of a time step in the calendar — month, quarter, year, holiday proximity. Unlike SARIMA's seasonal differencing, calendar dummies give tree-based models explicit information about which periods are structurally different.

### Key feature types:

- **Month dummies:** 11 binary columns (`month_2` through `month_12`); January is the reference category
- **Quarter dummies:** 3 binary columns (Q2–Q4)
- **Trend term:**  $t = 1, 2, \dots, T$  for linear time trend
- **Structural break indicator:** 0/1 for pre/post a known break (recession, pandemic, policy change)

## Construction in pandas:

```
df.index.month → integers 1–12
```

```
pd.get_dummies(df['month'], prefix='month',  
drop_first=True)
```

## Interpretation (Random Forest):

Month	Avg. Effect	Business Meaning
January (ref)	—	Post-holiday low
month_11	+1,800	Nov. ramp-up
month_12	+3,200	Holiday peak
month_7	+900	Summer surge
month_2	-400	Feb. dip

### Both encode seasonality, but differently:

- $\text{lag\_12} = y_{t-12}$ : level effect — captures the same magnitude as last year
- $\text{month\_12}$ : indicator effect — always adds the same fixed amount for December, regardless of last year's level

Use both when trend is changing.

*Socractic: A SARIMA model with  $s = 12$  handles seasonality through seasonal differencing. Why might month dummies improve a Random Forest but add no value to SARIMA?*

## Calendar features help when:

- The seasonal pattern is *fixed* across years (e.g., holiday retail, fiscal quarters)
- The model does not already capture seasonality via lag features (e.g., shallow trees, linear models without lag 12)
- Structural breaks are known and datable (add a 0/1 indicator for pre/post)

## Calendar features can hurt when:

- Seasonal patterns are *evolving* over time — the fixed-effect assumption breaks down
- `lag_12` already accounts for seasonality — month dummies add multicollinearity
- The dataset is small ( $n < 100$ ): 11 month dummies consume degrees of freedom faster than they add signal

Never use calendar dummies as the sole source of seasonality in a neural network (LSTM). LSTMs learn temporal patterns from the sequence itself; adding 11 overlapping dummies adds noise. Use them as supplementary features only when lag features are already present.

---

## Interaction & Ratio Features

Encoding change, not level

---

---

**Interaction and ratio features** encode *change* rather than *level*. They help tree-based models detect acceleration or deceleration in a series — a concept that requires computing ratios of two variables, something a single tree split cannot express directly.

### Key features:

- **Year-over-year (YoY) change:**  $\Delta^{(12)}y_t = y_{t-1}/y_{t-13} - 1$
- **Month-over-month (MoM) change:**  $\Delta^{(1)}y_t = y_{t-1}/y_{t-2} - 1$
- **Lag interaction:**  $y_{t-1} \times \mathbf{1}[\text{month} = 12]$  — lets the model use different slopes for December

## Year-over-year rate of change:

$$\text{YoY}_t = \frac{y_{t-1}}{y_{t-13}} - 1$$

Note: both  $y_{t-1}$  and  $y_{t-13}$  are lagged to prevent leakage at prediction time for target  $y_t$ .

## Month-over-month rate of change:

$$\text{MoM}_t = \frac{y_{t-1}}{y_{t-2}} - 1$$

**Why ratios instead of differences?** A 1,000-unit change means something different when the baseline is 10,000 vs. 100,000. Ratios are scale-free.

For RSXFS retail sales:

- $\text{YoY}_t > 0$ : economy is stronger than same month last year
- $\text{YoY}_t < 0$ : contraction signal
- $|\text{MoM}_t|$  large: abnormal month-to-month swing (data revision or shock)

These features help the model distinguish a Christmas peak from an unexpected demand shock.

*Socratic: YoY removes the seasonal level but still contains trend. What additional transformation would make  $\text{YoY}_t$  a stationary series?*

## Full feature engineering pipeline:

```
def make_features_extended(df,
                           lags=range(1, 13),
                           roll_windows=[3, 6, 12]):
    X = df[['value']].copy()
    # Lag features (ACF/PACF-guided, lags 1--12)
    for k in lags:
        X[f'lag_{k}'] = X['value'].shift(k)
    # Rolling stats (shift first to avoid leakage)
    for w in roll_windows:
        r = X['value'].shift(1).rolling(w)
        X[f'roll_mean_{w}'] = r.mean()
        X[f'roll_std_{w}'] = r.std()
        X[f'roll_min_{w}'] = r.min()
        X[f'roll_max_{w}'] = r.max()
    # Exponentially weighted mean
    X['ewm_alpha03'] = (
        X['value'].shift(1)
        .ewm(alpha=0.3, adjust=False).mean())
    # Calendar dummies (Jan = reference category)
    months = pd.Series(
        X.index.month, index=X.index)
    dums = pd.get_dummies(
        months, prefix='month',
        drop_first=True)
    X = pd.concat([X, dums], axis=1)
    # Return 36-column feature matrix
    return X.drop(columns='value').dropna()
```

---

## Feature Selection

Choosing the best subset of 36 candidate features

---

---

More features are not always better. With 36 features and  $n \approx 300$  observations, we risk overfitting — especially in linear models. **Feature selection** identifies the subset of features that maximizes out-of-sample predictive accuracy.

### Three strategies covered today:

1. **LASSO** (from Lecture 08):  $\ell_1$  penalty shrinks irrelevant feature coefficients to exactly zero. Implicit selection via the regularization path.
2. **Permutation importance**: measure how much test RMSE rises when a feature's values are randomly shuffled. Model-agnostic
3. **RFE CV**: Recursive Feature Elimination with cross-validation. Fit model, remove weakest feature, repeat until CV score stops improving

---

## LASSO regularization path (Lecture 08 connection):

$$\hat{\beta}^{\text{LASSO}} = \arg \min_{\beta} \sum_t \underbrace{(y_t - \mathbf{x}_t^\top \beta)^2}_{\text{prediction error}} + \lambda \underbrace{\|\beta\|_1}_{\text{selection penalty}}$$

As  $\lambda$  increases along the regularization path:

- **Lag 12 and lag 1** coefficients survive to large  $\lambda$  — high marginal importance
- **Rolling std features** shrink toward zero early — lower marginal importance
- At CV-optimal  $\lambda^*$ : many coefficients are exactly zero (automatic feature selection) (Tibshirani 1996)

*LASSO requires standardized features — all inputs should be on comparable scales (use StandardScaler inside a Pipeline to prevent leakage).*

1. Fit Random Forest on training set
2. Record **baseline** validation RMSE:  $\text{RMSE}_0$
3. For each feature  $j$ : shuffle column  $j$ , re-predict, record  $\text{RMSE}_j$
4.  $\text{Importance}_j = \text{RMSE}_j - \text{RMSE}_0$  (larger = more important)

### Advantages over split-count (impurity) importance:

- Evaluated on *validation* set — not biased toward high-cardinality features
- *Model-agnostic*: applies to LSTM, XGBoost, linear models
- Measures *marginal* contribution with all other features present

On RSXFS: *lag\_1*, *lag\_12*, *roll\_mean\_12*, *ewm\_alpha03* are the top four. See the importance preview in Section 2 (Rolling Statistics).

**Input:** estimator, feature matrix  $X$ , cross-validator

**Repeat:**

1. Fit estimator on all remaining features
2. Rank features by importance
3. Eliminate the lowest-ranked feature

**Select:** the subset size that maximizes mean CV score  
(Guyon and Elisseeff 2003)

**Key parameters:**

- step=1: eliminate one feature per round
- cv=TimeSeriesSplit(gap=1): prevents temporal leakage during selection
- min\_features\_to\_select=5: floor to avoid trivial models

```
from sklearn.feature_selection import RFEcv
from sklearn.ensemble import (
    RandomForestRegressor)
from sklearn.model_selection import (
    TimeSeriesSplit)

estimator = RandomForestRegressor(
    n_estimators=200, random_state=42)

tscv = TimeSeriesSplit(n_splits=5, gap=1)

selector = RFEcv(
    estimator=estimator,
    step=1,
    cv=tscv,
    scoring='neg_mean_squared_error',
    min_features_to_select=5)

selector.fit(X_tr, y_tr)

# Selected feature names
selected = X_tr.columns[
    selector.support_].tolist()
print(f"Selected: {len(selected)} feats")
```

---

## Pipeline Design

Preventing leakage in cross-validation

---

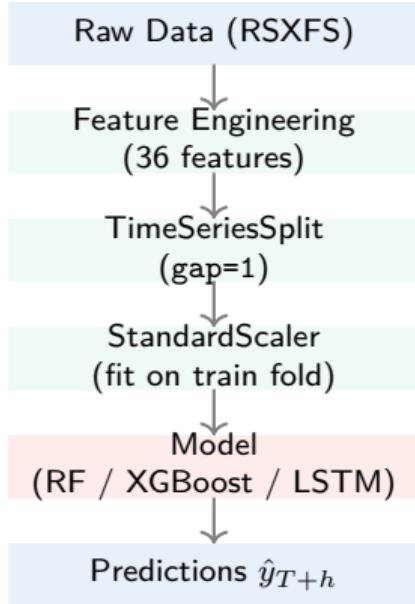
---

A **sklearn Pipeline** chains transformations and a model into a single object. Inside cross-validation, the pipeline re-fits every transformation on the training fold only — preventing the subtle leakage that occurs when scaling with statistics from the full dataset.

### Why pipelines matter for time-series CV:

- StandardScaler fitted on all data leaks test-set mean and variance into training
- With Pipeline + TimeSeriesSplit: scaler is fitted fresh on each train fold
- Same principle applies to RFECV and any imputer — fit only on train, transform both train and test (Pedregosa et al. 2011)

*Rule of thumb: if a transformation looks at the target variable (e.g., LabelEncoder, WoE encoding) or at the distribution of X (e.g., StandardScaler, PCA), it must go inside the pipeline.*



## What happens inside each CV fold:

1. **Split:** `TimeSeriesSplit(gap=1)` — 1-step gap prevents leakage
2. **Fit scaler:** `StandardScaler` on train fold *only*
3. **Transform:** same params applied to train and val
4. **Fit model:** on scaled train fold
5. **Score:** RMSE on scaled val fold
6. **Report:** average across all folds

Fitting `StandardScaler` before CV on all data: validation mean/std seep into training. Typical RMSE improvement from fixing this: 50–200 units on RSXFS.

## Step 1: Construct the pipeline

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import (
    StandardScaler)
from sklearn.ensemble import (
    RandomForestRegressor)
from sklearn.model_selection import (
    TimeSeriesSplit, cross_val_score)
import numpy as np

# Build pipeline (order matters!)
pipe = Pipeline([
    ('scaler', StandardScaler()),
    ('model', RandomForestRegressor(
        n_estimators=300,
        random_state=42))])

tscv = TimeSeriesSplit(
    n_splits=5, gap=1)
```

## Step 2: Cross-validate and predict

```
# CV: scaler fit on each train fold
scores = cross_val_score(
    pipe,
    X.values,
    y.values,
    cv=tscv,
    scoring=(
        'neg_mean_squared_error'))
rmse_cv = np.sqrt(-scores.mean())
print(f"CV RMSE: {rmse_cv:.0f}")

# Refit on full training data
pipe.fit(X_tr.values, y_tr.values)

# Predict on held-out test set
y_pred = pipe.predict(
    X_te.values)
rmse_te = np.sqrt(
    np.mean((y_te - y_pred)**2))
print(f"Test RMSE: {rmse_te:.0f}")
```

---

## Application to Forecasting

Baseline vs. extended features on RSXFS retail sales

---

---

Apply `make_features_extended()` (36 features) to RSXFS retail sales and compare against the baseline 26-feature set from Lectures 07–10. Measure the marginal contribution of better features holding the model class fixed.

### Evaluation design:

- **Baseline** (26 features): 12 lags, 6 rolling means, 6 rolling stds, 2 ratios — as in Lectures 07–10
- **Extended** (36 features): adds rolling min/max, EWM, and 11 month dummies
- **Models**: Elastic Net, Random Forest, XGBoost, LSTM (median over 5 seeds)
- **Holdout**: same chronological 15% test set used throughout the course

## Test-set RMSE on RSXFS:

Model	26f	36f	$\Delta$ RMSE
Seasonal Naïve	4 210	—	—
SARIMA(1,1,1)(1,1,1) <sub>12</sub>	2 840	—	—
Elastic Net	2 540	2 410	-130
Random Forest	2 380	2 210	-170
XGBoost	2 250	2 050	-200
LSTM (2-layer, $T=24$ )	2 180	1 920	-260

Values are illustrative. LSTM = median over 5 random seeds.

## Interpretation:

- All four ML models improve with richer features; LSTM gains most
- The improvement is larger for more flexible models — extra features give trees more split choices
- Elastic Net gains least: its penalty already prevents overfitting to noisy features

Feature engineering delivers larger improvements than switching model class. RMSE:  
**LSTM 36f < XGBoost 36f < LSTM 26f**:  
the feature gap dominates the model gap.

---

## Takeaways and References

What we learned and where to go next

---

**Lag features** convert time series to supervised regression; use ACF and PACF to select which lags to include (Box et al. 2015).

**Rolling statistics** (mean, std, min, max, EWM) capture local trend and volatility at multiple scales. Always `shift(1)` before rolling to prevent leakage.

**Calendar dummies** help tree models detect recurring peaks and troughs; they add less value when `lag_12` is already present.

**Feature selection** (LASSO, permutation importance, RFECV) reduces overfitting and identifies the most informative signals (Guyon and Elisseeff 2003; Molnar 2022).

**Pipelines** are essential for leakage-free cross-validation. Fit all transformations inside the CV loop, not on the full dataset (Pedregosa et al. 2011).

**Feature engineering beats model selection:** on RSXFS, moving from 26 to 36 features reduced RMSE by 130–260 units across all model classes.

**Preview of Lecture 12:** Capstone and Applications — combining the full Lectures 01–11 toolkit on a business case study with model selection, uncertainty quantification, and presentation-ready visualizations.

- 
-  Box, George E. P. et al. (2015). *Time Series Analysis: Forecasting and Control*. 5th. Hoboken, NJ: Wiley.
  -  Guyon, Isabelle and André Elisseeff (2003). "An Introduction to Variable and Feature Selection". In: *Journal of Machine Learning Research* 3, pp. 1157–1182.
  -  Molnar, Christoph (2022). *Interpretable Machine Learning: A Guide for Making Black Box Models Explainable*. 2nd. Open access. URL:  
<https://christophm.github.io/interpretable-ml-book/>.
  -  Pedregosa, Fabian et al. (2011). "Scikit-learn: Machine Learning in Python". In: *Journal of Machine Learning Research* 12, pp. 2825–2830.
  -  Tibshirani, Robert (1996). "Regression Shrinkage and Selection via the Lasso". In: *Journal of the Royal Statistical Society: Series B* 58.1, pp. 267–288.